

Algorithms Lab

September 26, 2012

1 C++ STL

2 Exercise hints

We assume that you ...

- know how to code
- know the basic algorithms
- are able to turn a reasonably detailed idea into a working code! (otherwise the hints we give will not be very useful)
- are familiar with the \mathcal{O} notation and the idea of estimating the runtime of your program based on the input size

1 C++ STL

2 Exercise hints

- Coding data structure from scratch - painful and error prone
- STL - *map*, *set*, *queue*, *priority queue*
 - map - $\mathcal{O}(\log n)$ lookups
 - set - $\mathcal{O}(\log n)$ set membership tracking
 - queue - FIFO
 - priority queue - $\mathcal{O}(\log n)$ retrieval of an element with the highest priority, insertion
- look at the documentation

Sorting

```
1 #include <algorithm>
2 ...
3 vector<int> a;
4 ...
5 sort(a.begin(), a.end());
```

Sorting

```
1  #include <algorithm>
2  ...
3  struct Data {
4      int key;
5
6      bool operator<(const Data &a) const {
7          return key < a.key;
8      }
9  };
10 ...
11 vector<Data> a;
12 ...
13 sort(a.begin(), a.end());
```

Same applies for set, map and priority queue

```
1  #include <set>
2  ...
3  struct Data {
4      int key;
5
6      bool operator<(const Data &a) const {
7          return key < a.key;
8      }
9  };
10 ...
11 set<Data> a;
12 ...
13 set<Data>::iterator itr = a.find(data);
```


1 C++ STL

2 Exercise hints

Exercise 1 - False coin

Two approaches

- (slower) Guess a solution
 - assume a coin is heavier (or lighter)
 - check consistency with the input
 - exactly one such coin – we found a solution
- (faster) Keep the track of possibly false coins

Exercise 1 - False coin

Two approaches

- (slower) Guess a solution
 - assume a coin is heavier (or lighter)
 - check consistency with the input
 - exactly one such coin – we found a solution
- (faster) Keep the track of possibly false coins

Exercise 1 - False coin

Two approaches

- (slower) Guess a solution
 - assume a coin is heavier (or lighter)
 - check consistency with the input
 - exactly one such coin – we found a solution
- (faster) Keep the track of possibly false coins

Exercise 1 - False coin

Two approaches

- (slower) Guess a solution
 - assume a coin is heavier (or lighter)
 - check consistency with the input
 - exactly one such coin – we found a solution
- (faster) Keep the track of possibly false coins

Exercise 1 - False coin

Two approaches

- (slower) Guess a solution
 - assume a coin is heavier (or lighter)
 - check consistency with the input
 - exactly one such coin – we found a solution
- (faster) Keep the track of possibly false coins

Exercise 1 - False coin

Two approaches

- (slower) Guess a solution
 - assume a coin is heavier (or lighter)
 - check consistency with the input
 - exactly one such coin – we found a solution
- (faster) Keep the track of possibly false coins

Exercise 1 - False coin - Summary

Trick learned

- guess a solution
- guess an interval in which solution lies, properties of a solution, etc.

Exercise 2 - Race Tracks

- Smallest number of hops
- Recall – BFS finds the smallest number of steps from vertex v_1 to v_2
- Model the input as a graph

Exercise 2 - Race Tracks

Naive model

- vertices - points in the grid
- edge between two vertices – hop from one point to the other
- hop depends on our current velocity
- different velocities give different hops from the same vertex
- graph with edges which dynamically change? doesn't sound good ...

Exercise 2 - Race Tracks

Naive model

- vertices - points in the grid
- edge between two vertices – hop from one point to the other
- hop depends on our current velocity
- different velocities give different hops from the same vertex
- graph with edges which dynamically change? doesn't sound good ...

Exercise 2 - Race Tracks

Better model

- vertex is a pair (p, v) – p is a point in the grid, v is the velocity which we had upon arriving to that point
 - Start vertex – $(s, (0, 0))$
 - End vertices – $(t, (i, j))$, for all $-3 \leq i, j \leq 3$
- Number of vertices – at most $(30 \times 30) \times (7 \times 7) \approx 45000$
- Degree of each vertex at most 9 – at most 405000 edges
- Edges can be deduced from the "label" of a vertex – no need to store the whole graph explicitly

Exercise 2 - Race Tracks

Better model

- vertex is a pair (p, v) – p is a point in the grid, v is the velocity which we had upon arriving to that point
 - Start vertex – $(s, (0, 0))$
 - End vertices – $(t, (i, j))$, for all $-3 \leq i, j \leq 3$
- Number of vertices – at most $(30 \times 30) \times (7 \times 7) \approx 45000$
- Degree of each vertex at most 9 – at most 405000 edges
- Edges can be deduced from the "label" of a vertex – no need to store the whole graph explicitly

Exercise 2 - Race Tracks

Better model

- vertex is a pair (p, v) – p is a point in the grid, v is the velocity which we had upon arriving to that point
 - Start vertex – $(s, (0, 0))$
 - End vertices – $(t, (i, j))$, for all $-3 \leq i, j \leq 3$
- Number of vertices – at most $(30 \times 30) \times (7 \times 7) \approx 45000$
- Degree of each vertex at most 9 – at most 405000 edges
- Edges can be deduced from the "label" of a vertex – no need to store the whole graph explicitly

Exercise 2 - Race Tracks - Summary

Prerequisites

- BFS!

Trick learned

- Vertices of a graph can be represented with more complex objects than just numbers from 1 to n
- No need to store the graph explicitly in order to perform a BFS or similar algorithm

Exercise 3 - Forced Landing - First approach

Observations

- Equivalent – given an $h \times w$ matrix with 0/1 elements, find an $r \times s$ submatrix with all 0 elements
 - up to a "constant" – we need a $(r - 1) \times (s - 1)$ submatrix, etc.
- r and s are fixed – top-left corner completely describes the rectangle

Exercise 3 - Forced Landing - First approach

Observations

- Equivalent – given an $h \times w$ matrix with 0/1 elements, find an $r \times s$ submatrix with all 0 elements
 - up to a "constant" – we need a $(r - 1) \times (s - 1)$ submatrix, etc.
- r and s are fixed – top-left corner completely describes the rectangle

Exercise 3 - Forced Landing - First approach

Observations

- Equivalent – given an $h \times w$ matrix with 0/1 elements, find an $r \times s$ submatrix with all 0 elements
 - up to a "constant" – we need a $(r - 1) \times (s - 1)$ submatrix, etc.
- r and s are fixed – top-left corner completely describes the rectangle

Exercise 3 - Forced Landing - First approach

Easy solution

- try all possible places for the top-left corner
- check whether the induced rectangle is empty (has all 0 elements)
- complexity
 - $\mathcal{O}(h \cdot w)$ possible rectangles
 - checking whether it is empty by simply iterating through all its elements – $\mathcal{O}(r \cdot s)$
 - in total $\mathcal{O}(hwrs)$ – too slow

Exercise 3 - Forced Landing - First approach

Easy solution

- try all possible places for the top-left corner
- check whether the induced rectangle is empty (has all 0 elements)
- complexity
 - $\mathcal{O}(h \cdot w)$ possible rectangles
 - checking whether it is empty by simply iterating through all its elements – $\mathcal{O}(r \cdot s)$
 - in total $\mathcal{O}(hws)$ – too slow

Exercise 3 - Forced Landing - First approach

- try to solve a simpler problem first – suppose that $h = r = 1$
- an array instead of a matrix

Prefix sums

Given an array A , set $s[i] := \sum_{j=1}^i A[j]$. Additionally, set $A[0] := 0$.

- subarray $A[i], \dots, A[i + s - 1]$ is empty \leftrightarrow
 $A[i + s - 1] - A[i - 1] = 0$
- apply similar trick to a matrix
 - reduces the complexity of checking whether a rectangle is empty from $\mathcal{O}(rs)$ to $\mathcal{O}(1)$!
 - in total $\mathcal{O}(hw)$ – works well for small and medium testset

Exercise 3 - Forced Landing - Second approach

Idea – move two vertical lines at distance s from left to right (*scanlines*)

Observation 1

- Suppose that the left scanline is at some position x_ℓ
- let X_ℓ denote the set of all trees with x coordinate in the interval $[x_\ell, x_\ell + s]$
- if there exist two trees in X_ℓ whose y coordinate differ by at least r , and no other tree in X_ℓ has y coordinate between these two – possible landing rectangle!

Exercise 3 - Forced Landing - Second approach

Idea – move two vertical lines at distance s from left to right (*scanlines*)

Observation 1

- Suppose that the left scanline is at some position x_ℓ
- let X_ℓ denote the set of all trees with x coordinate in the interval $[x_\ell, x_\ell + s]$
- if there exist two trees in X_ℓ whose y coordinate differ by at least r , and no other tree in X_ℓ has y coordinate between these two – possible landing rectangle!

Observation 2

- we cannot move scanlines continuously ...
- what are the points of interest of either left or the right scanline?
 - hint – trees

For a given positions of scanlines, how to compute whether there exist two trees from the observation 1?

Observation 2

- we cannot move scanlines continuously ...
- what are the points of interest of either left or the right scanline?
 - hint – trees

For a given positions of scanlines, how to compute whether there exist two trees from the observation 1?

Observation 2

- we cannot move scanlines continuously ...
- what are the points of interest of either left or the right scanline?
 - hint – trees

For a given positions of scanlines, how to compute whether there exist two trees from the observation 1?

Observation 2

- we cannot move scanlines continuously ...
- what are the points of interest of either left or the right scanline?
 - hint – trees

For a given positions of scanlines, how to compute whether there exist two trees from the observation 1?

Exercise 3 - Forced Landing - Second approach

Note

Filling out the missing details, and especially implementing it correctly, is very difficult! However, solving it is very rewarding!

Exercise 3 - Forced Landing - Summary

Tricks learned

- Prefix sums
- Scanline

General hint

- If a problem is modeled in k dimension, and it seems difficult, try first solving it if it has one less dimension and then extend the idea.
- *Note: this is a hint, not a rule! (so it might not always work)*

Exercise 3 - Forced Landing - Summary

Tricks learned

- Prefix sums
- Scanline

General hint

- If a problem is modeled in k dimension, and it seems difficult, try first solving it if it has one less dimension and then extend the idea.
- *Note:* this is a hint, not a rule! (so it might not always work)

Exercise 4 - Demolition

Easy solution – blow up floors $1, \dots, \ell$

- blow up ℓ -th floor and simulate what happens after that
- complexity – $\mathcal{O}(n)$
- iterate over all floors – $\mathcal{O}(n^2)$ in total

Exercise 4 - Demolition

Easy solution – blow up floors $1, \dots, \ell$

- blow up ℓ -th floor and simulate what happens after that
- complexity – $\mathcal{O}(n)$
- iterate over all floors – $\mathcal{O}(n^2)$ in total

Exercise 4 - Demolition

Easy solution – blow up floors $1, \dots, \ell$

- blow up ℓ -th floor and simulate what happens after that
- complexity – $\mathcal{O}(n)$
- iterate over all floors – $\mathcal{O}(n^2)$ in total

Exercise 4 - Demolition

- S_ℓ – set of floors which you need to blow up additionally if the highest blown up floor is ℓ
- Relation of S_ℓ and $S_{\ell+1}$?
- Speed-up the simulation
 - using prefix sums – check in $\mathcal{O}(1)$ whether a floor i collapses if floors $i+1, \dots, j$ collapse
 - for each floor in S_ℓ – check whether it survives the additional weight of the floor $\ell+1$
 - however – still $\mathcal{O}(n^2)$... but enough for 50 points!

Exercise 4 - Demolition

- S_ℓ – set of floors which you need to blow up additionally if the highest blown up floor is ℓ
- Relation of S_ℓ and $S_{\ell+1}$?
- Speed-up the simulation
 - using prefix sums – check in $\mathcal{O}(1)$ whether a floor i collapses if floors $i+1, \dots, j$ collapse
 - for each floor in S_ℓ – check whether it survives the additional weight of the floor $\ell+1$
 - however – still $\mathcal{O}(n^2)$... but enough for 50 points!

Exercise 4 - Demolition

- S_ℓ – set of floors which you need to blow up additionally if the highest blown up floor is ℓ
- Relation of S_ℓ and $S_{\ell+1}$?
- Speed-up the simulation
 - using prefix sums – check in $\mathcal{O}(1)$ whether a floor i collapses if floors $i+1, \dots, j$ collapse
 - for each floor in S_ℓ – check whether it survives the additional weight of the floor $\ell+1$
 - however – still $\mathcal{O}(n^2)$... but enough for 50 points!

Exercise 4 - Demolition

Getting to $\mathcal{O}(n \log n)$

- consider some $f_1, f_2 \in S_\ell$
- define $A(f, \ell) := \text{cap}[f] - \sum_{i=1}^{\ell} w[i]$
- $A(f_1, \ell) > A(f_2, \ell)$
 - if f_2 doesn't collapse, neither will f_1
- simulation for the level $\ell + 1$ (hint)
 - get the element $f \in S_\ell$ with the smallest $A(f, \ell)$
 - check whether it collapses ...
- $A(f_1, \ell) > A(f_2, \ell) \rightarrow A(f_1, \ell') > A(f_2, \ell')$ for every $\ell' > \ell$

Exercise 4 - Demolition

Getting to $\mathcal{O}(n \log n)$

- consider some $f_1, f_2 \in S_\ell$
- define $A(f, \ell) := \text{cap}[f] - \sum_{i=1}^{\ell} w[i]$
- $A(f_1, \ell) > A(f_2, \ell)$
 - if f_2 doesn't collapse, neither will f_1
- simulation for the level $\ell + 1$ (hint)
 - get the element $f \in S_\ell$ with the smallest $A(f, \ell)$
 - check whether it collapses ...
- $A(f_1, \ell) > A(f_2, \ell) \rightarrow A(f_1, \ell') > A(f_2, \ell')$ for every $\ell' > \ell$

General hints

- Don't rush to find the fastest solution
- Don't expect to come up with the brilliant idea just by looking at the problem – work towards it!
- Not everyone can solve every problem for 100 points; however, we will usually give a non-negligible number of points for slower solutions
 - if the problem requires $\mathcal{O}(n \log n)$ solution or faster, coding a $\mathcal{O}(n^2)$ one will very likely give you a certain number of points – which might save you on the exam!
 - not every slower solution will earn you some points – $\mathcal{O}(2^n)$ solution, for example, might be too slow if the optimal solution is $\mathcal{O}(n)$

General hints

- Don't rush to find the fastest solution
- Don't expect to come up with the brilliant idea just by looking at the problem – work towards it!
- Not everyone can solve every problem for 100 points; however, we will usually give a non-negligible number of points for slower solutions
 - if the problem requires $\mathcal{O}(n \log n)$ solution or faster, coding a $\mathcal{O}(n^2)$ one will very likely give you a certain number of points – which might save you on the exam!
 - not every slower solution will earn you some points – $\mathcal{O}(2^n)$ solution, for example, might be too slow if the optimal solution is $\mathcal{O}(n)$

General hints

- Don't rush to find the fastest solution
- Don't expect to come up with the brilliant idea just by looking at the problem – work towards it!
- Not everyone can solve every problem for 100 points; however, we will usually give a non-negligible number of points for slower solutions
 - if the problem requires $\mathcal{O}(n \log n)$ solution or faster, coding a $\mathcal{O}(n^2)$ one will very likely give you a certain number of points – which might save you on the exam!
 - not every slower solution will earn you some points – $\mathcal{O}(2^n)$ solution, for example, might be too slow if the optimal solution is $\mathcal{O}(n)$

General hints

- Don't rush to find the fastest solution
- Don't expect to come up with the brilliant idea just by looking at the problem – work towards it!
- Not everyone can solve every problem for 100 points; however, we will usually give a non-negligible number of points for slower solutions
 - if the problem requires $\mathcal{O}(n \log n)$ solution or faster, coding a $\mathcal{O}(n^2)$ one will very likely give you a certain number of points – which might save you on the exam!
 - not every slower solution will earn you some points – $\mathcal{O}(2^n)$ solution, for example, might be too slow if the optimal solution is $\mathcal{O}(n)$

More fun interactively.

Remember

- Test locally!
- Times are deterministic. Submitting several times is a waste of resources

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit taking too long

compiler-error fix it, compile locally

segmentation-fault SIGSEGV: memory screwup

assertion-failure SIGABRT: memory screwup or `assert()`
or uncaught exception

run-error nonzero exit status

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer **fix it, make testsets**

timelimit taking too long

compiler-error fix it, compile locally

segmentation-fault SIGSEGV: memory screwup

assertion-failure SIGABRT: memory screwup or `assert()`
or uncaught exception

run-error nonzero exit status

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit rethink approach

compiler-error fix it, compile locally

segmentation-fault SIGSEGV: memory screwup

assertion-failure SIGABRT: memory screwup or `assert()`
or uncaught exception

run-error nonzero exit status

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit rethink approach

compiler-error **fix it, compile locally**

segmentation-fault SIGSEGV: memory screwup

assertion-failure SIGABRT: memory screwup or `assert()`
or uncaught exception

run-error nonzero exit status

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit rethink approach

compiler-error fix it, compile locally

segmentation-fault use valgrind

assertion-failure SIGABRT: memory screwup or `assert()`
or uncaught exception

run-error nonzero exit status

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit rethink approach

compiler-error fix it, compile locally

segmentation-fault use valgrind

assertion-failure use gdb or valgrind

run-error nonzero exit status

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit rethink approach

compiler-error fix it, compile locally

segmentation-fault use valgrind

assertion-failure use gdb or valgrind

run-error **nonzero exit status**

forbidden bad syscall or other safety limit

presentation-error whitespace differences

Judge Results

correct you win!

wrong-answer fix it, make testsets

timelimit rethink approach

compiler-error fix it, compile locally

segmentation-fault use valgrind


assertion-failure use gdb or valgrind

run-error nonzero exit status


forbidden **stop it**

presentation-error whitespace differences

Distribution of results



Distribution of results



Debugging utilities discussed

valgrind

- + easy to use
- + verbose
- occasionally hard to interpret
- no interactive investigation
- issues with rounding in CGAL

gdb

- + interactive
- + many possibilities
- harder to use
- post-segfault is often too late

debugging prints

GDB cheatsheet

<code>gdb prog</code>	start gdb on the <i>prog</i>
<code>run args</code>	run program with <i>args</i> as in shell
<code>bt</code>	print stack trace (backtrace)
<code>bt -N</code>	— of outermost <i>N</i> frames
<code>up</code>	up one frame (towards <code>main</code>)
<code>down</code>	down one frame
<code>print expr</code>	print the value of <i>expr</i>
<code>break func</code>	add breakpoint when entering <i>func</i>
<code>break line</code>	add breakpoint before <i>line</i>
<code>continue</code>	continue execution
<code>next</code>	— until next line in this function
<code>step</code>	— until next line anywhere
<code>quit</code>	quit

Take-aways of interactive session

- 01 Use assertions to catch input formatting mistakes (useful with your own testsets!)
- 02 Use `-Wall` at all times
- 03 `switch()` is finicky
- 04 Use `valgrind`
Beware of uninitialized values
- 05 Beware of array overruns; try `.at(i)`
- 06 Pointers invalidated across vector resizing
- 07 Be clear on static/heap/automatic data model
- 08 Stack overflow is also a segfault
- 09 Violating contracts can break STL algos